

AI & ML for Data Scientists

Class 8: Classification and Risk Quant

Lei Ge (葛雷)

Quant RUC (人大量化)

November 5, 2025



Introduction to Classification

Logistic Regression

Performance Measures

Multiclass Classification

Introduction to Classification

Logistic Regression

Performance Measures

Multiclass Classification

What is Classification?

- Classification is the task of predicting a label (class).
- Binary classification: two classes (e.g., spam vs. ham).
- Multiclass classification: more than two classes.
- Multilabel and multioutput classifications

Applications of Classification Models

Finance:

- Credit scoring: Predict if a customer will default on a loan.
- Fraud detection: Classify whether a transaction is fraudulent.

Healthcare:

- Disease diagnosis: Predict presence of conditions (e.g., cancer, diabetes).
- Risk stratification: Identify high-risk patients.

Marketing:

- Customer segmentation.
- Predict likelihood to respond to a campaign.

LLM is also a Classification Model

- Why ?
- Word (token) prediction from the huge dictionary

Related Jobs: Risk Quants

- Credit Risk Quant: Models borrower defaults
- Market Risk Quant: Focuses on pricing derivatives and hedging strategies risk
- Operational Risk Quant: Analyzes tail-risk events (e.g., fraud, system failures)
- Model Risk Quant: model validation to prevent risks from modeling

Why Linear Regression not for classification

- Why Linear Regression not for classification?
- Let me give you a example

Sigmoid function for the non-linearity

Equation 4-14. Logistic function

$$\sigma(t) = \frac{1}{1 + \exp(-t)}$$




Figure 4-21. Logistic function

What is Logistic Regression?

- A statistical model for binary classification
- Predicts the probability of an event occurring
- Output ranges between 0 and 1
- Uses the logistic (sigmoid) function:

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

Logistic Regression

- Traditional Logistic Regression:

$$\hat{P} = \sigma(x\theta^T) = \frac{1}{1 + e^{-x\theta^T}}$$

- Other machine learning models:

$$\hat{P} = \sigma(h_{\theta}(x)) = \frac{1}{1 + e^{-h_{\theta}(x)}}$$

- $\sigma(\cdot)$ is also called activation function (create non-linearity)
- $h_{\theta}(x)$ is also called score function, it can be ANN, Xgboost, random forest ...

From Soft Output ($\hat{P}$) to Hard Output ($\hat{Y}$)

Hard Output:

$$\hat{Y} = \begin{cases} 1, & \text{if } \hat{P} \geq \tau \\ 0, & \text{if } \hat{P} < \tau \end{cases} \quad \text{where } \tau \in [0, 1]$$

- $\hat{P}$ — soft output (probabilistic prediction)
- $\hat{Y}$ — hard output (binary class decision)
- τ — threshold (commonly set to 0.5)
- threshold low $\rightarrow$ more $\hat{Y}$ become 1
- Soft output $\hat{Y}$ is more important in finance (Why?)

Cross-entropy Loss Function (Cost function)

- Binary cross-entropy loss (log loss):

$$L(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log(\hat{P}^{(i)}) + (1 - y^{(i)}) \log(1 - \hat{P}^{(i)})]$$

- Where:
 - m = number of training examples
 - $y^{(i)}$ = true label (0 or 1)
 - $\hat{P}^{(i)}$ = model predicted the probability, also called **soft output**

Cross-entropy Loss?

- Why we use Cross-entropy Loss?
- Bernoulli Distribution + MLE
- Pencil and paper time!!!

Optimization: Gradient Descent

- Iterative optimization algorithm
- Update rule:

$$\theta_{t+1} = \theta_t - \alpha \frac{\partial L(\theta)}{\partial \theta_t}$$

- Partial derivative:

$$\theta_{t+1} = \theta_t - \alpha \frac{\partial L(\theta)}{\partial \hat{P}(i)} \frac{\partial \sigma(h_{\theta}(x^{(i)}))}{\partial \theta}$$

Example: Training a Binary Classifier

- Example: Classify whether a digit is 5 or not.
- Model: `SGDClassifier` from Scikit-Learn.
- Uses stochastic gradient descent.

Introduction to Classification

Logistic Regression

Performance Measures

Multiclass Classification

Measuring Performance

- Accuracy is not reliable for imbalanced datasets. (Why? Let me give a example)
- Confusion matrix is more informative.
- Precision and recall are better metrics.

Confusion Matrix

		Predicted Class		
		Positive	Negative	
Actual Class	Positive	True Positive (TP)	False Negative (FN) Type II Error	Sensitivity $\frac{TP}{(TP + FN)}$
	Negative	False Positive (FP) Type I Error	True Negative (TN)	Specificity $\frac{TN}{(TN + FP)}$
		Precision $\frac{TP}{(TP + FP)}$	Negative Predictive Value $\frac{TN}{(TN + FN)}$	Accuracy $\frac{TP + TN}{(TP + TN + FP + FN)}$

Confusion Matrix on textbook (labels different)




Figure 3-2. An illustrated confusion matrix shows examples of true negatives (top left), false positives (top right), false negatives (lower left), and true positives (lower right)

Precision vs Recall

- Precision = $\frac{TP}{TP+FP}$
- Recall = $\frac{TP}{TP+FN}$

Precision/Recall Trade-off

- You can adjust the decision threshold.
- $Precision(threshold)$ vs $Recall(threshold)$
- Precision is increasing in when threshold increase, Recall is decreasing when threshold is increasing.
- Higher precision $\Rightarrow$ lower recall and vice versa.

Precision/Recall Trade-off

Textbook (HandOn page 97): Predict Number 5 by change its threshold




Figure 3-3. Decision threshold and precision/recall tradeoff

Precision/Recall Trade-off




Figure 3-5. Precision versus recall

F1 Score

$$\begin{bmatrix} TN & FP \\ FN & TP \end{bmatrix}$$

- Precision = $\frac{TP}{TP+FP}$
- Recall = $\frac{TP}{TP+FN}$
- F1 Score = $F_1 = 2 \cdot \frac{Precision \cdot Recall}{Precision + Recall}$ s.t. Harmonic mean of precision and recall

ROC Curve

- ROC = Receiver Operating Characteristic curve.
- Plots True Positive Rate (Recall) against False Positive Rate.
- $TPR = \frac{TP}{TP+FN}$
- $FPR = \frac{FP}{FP+TN}$
- Helps evaluate classifier performance at all thresholds.

ROC Curve




Figure 3-6. This ROC curve plots the false positive rate against the true positive rate for all possible thresholds; the red circle highlights the chosen ratio (at 43.68% recall)

ROC Curve: Compare models




Figure 3-7. Comparing ROC curves: the Random Forest classifier is superior to the SGD classifier because its ROC curve is much closer to the top-left corner, and it has a greater AUC

Understanding AUC (Area Under Curve)

- AUC is the area under the ROC curve.
- $AUC = 1$: perfect classifier.
- $AUC = 0.5$: random guessing.
- The higher the AUC, the better the model distinguishes between classes.
- Useful when comparing multiple classifiers.

Introduction to Classification

Logistic Regression

Performance Measures

Multiclass Classification

Multiclass Classification: Softmax Function

Sigmoid vs Softmax

From Sigmoid → Softmax

Softmax Function

For a K -class classification problem:

$$\hat{p}_k = \frac{e^{z_k}}{\sum_{j=1}^K e^{z_j}} \quad \text{where } k = 1, \dots, K \quad (1)$$

Cross-Entropy Loss (Softmax Loss)

Given the softmax output $\hat{p}$ and the true label y , the loss for **one observation** is:

$$\mathcal{L} = - \sum_{k=1}^K y_k \log(\hat{p}_k)$$

- **K is the number of classes**
- Penalizes incorrect predictions more when the predicted probability is low
- Encourages the correct class to have a high predicted probability

Reference

1. Hands-on Machine Learning with Scikit-Learn, Keras and TensorFlow (3rd edition)
2. <https://encord.com/glossary/confusion-matrix/>
3. Kaggle
4. Wikipedia
5. ChatGPT
6. DeepSeek